

CSCI 4050 / 6050

Software Engineering

Introduction to

Relational Databases

and

Object-Relational Mapping



Outline

- Relational Databases
- Defining tables
- Structured Query Language
- Relational Database Management Systems
- JDBC
- Object-Relational Mapping
- Object persistence

Slides by Prof Kris Kouchut: In part, based on slides by Bill Howe at Portland State, and other sources



Introduction

- Database
 - a collection of persistent data
- Database Management System (DBMS)
 - a software system that supports creation, population, querying, and administering of a database
- Relational Database Management System
 - DBMS based on a Relational Model, created by Edgar Codd in 1969



Relational Database

- Relational Database (RDB)
 - Consists of a number of *tables* and a single *schema* (definition of tables and their attributes)
 - For example:
Student (sid, name, login, age, gpa)

Student identifies the table, while
sid, name, login, age, gpa identify attributes
sid is the primary key



An Example Table

- Student (*sid*: integer, *name*: string, *login*: string, *age*: integer, *gpa*: real)

| <u>sid</u> | name | login | age | gpa |
|------------|---------|---------------|-----|-----|
| 50000 | Dave | dave@cs | 19 | 3.3 |
| 53666 | Jones | jones@cs | 18 | 3.4 |
| 53688 | Smith | smith@ee | 18 | 3.2 |
| 53650 | Smith | smith@math | 19 | 3.8 |
| 53831 | Madayan | madayan@music | 11 | 1.8 |
| 53832 | Guldu | guldu@music | 12 | 2.0 |



An Example Table

Data types

- Student (*sid*: integer, *name*: string, *login*: string, *age*: integer, *gpa*: real)

| <u>sid</u> | name | login | age | gpa |
|------------|---------|---------------|-----|-----|
| 50000 | Dave | dave@cs | 19 | 3.3 |
| 53666 | Jones | jones@cs | 18 | 3.4 |
| 53688 | Smith | smith@ee | 18 | 3.2 |
| 53650 | Smith | smith@math | 19 | 3.8 |
| 53831 | Madayan | madayan@music | 11 | 1.8 |
| 53832 | Guldu | guldu@music | 12 | 2.0 |



An Example Table

- Student is a **relation** on $\text{Int} \times \text{String} \times \text{String} \times \text{Int} \times \text{Float}$
- **Row** represents a tuple (related values)

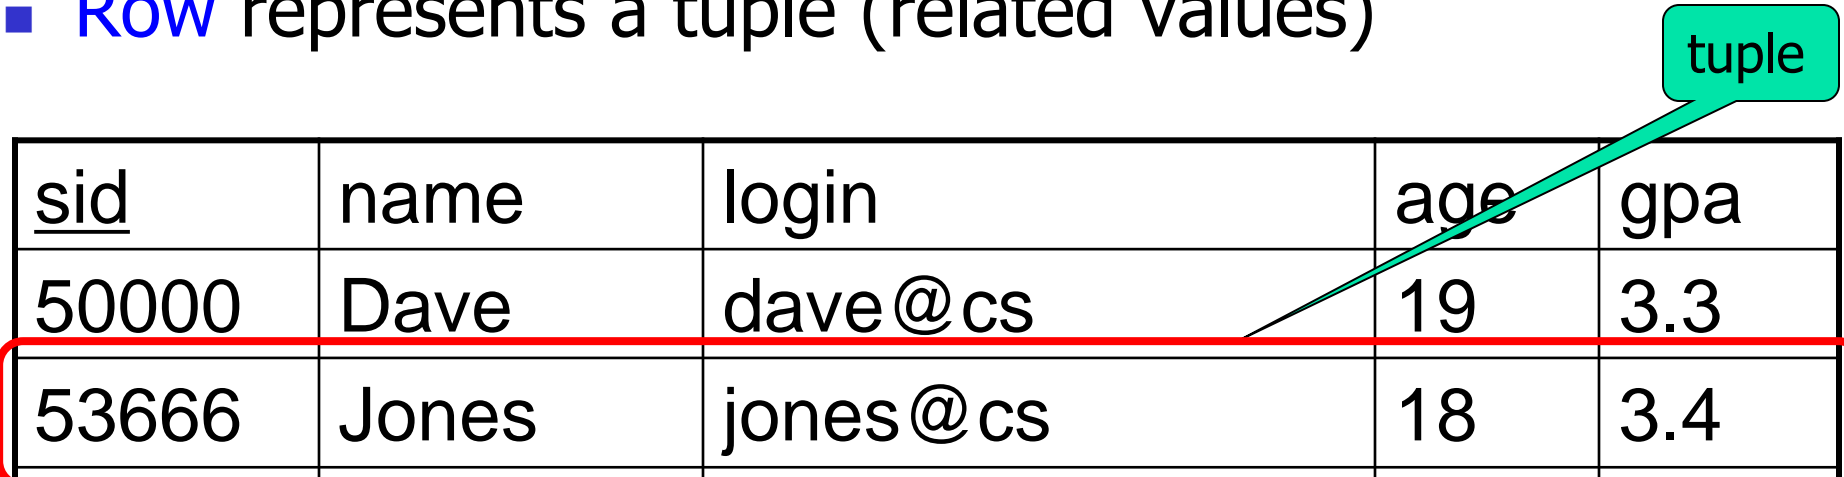
attributes

| <u>sid</u> | name | login | age | gpa |
|------------|---------|---------------|-----|-----|
| 50000 | Dave | dave@cs | 19 | 3.3 |
| 53666 | Jones | jones@cs | 18 | 3.4 |
| 53688 | Smith | smith@ee | 18 | 3.2 |
| 53650 | Smith | smith@math | 19 | 3.8 |
| 53831 | Madayan | madayan@music | 11 | 1.8 |
| 53832 | Guldu | guldu@music | 12 | 2.0 |



An Example Table

- Student is a **relation** on $\text{Int} \times \text{String} \times \text{String} \times \text{Int} \times \text{Float}$
- **Row** represents a tuple (related values)





| <u>sid</u> | name | login | age | gpa |
|------------|---------|---------------|-----|-----|
| 50000 | Dave | dave@cs | 19 | 3.3 |
| 53666 | Jones | jones@cs | 18 | 3.4 |
| 53688 | Smith | smith@ee | 18 | 3.2 |
| 53650 | Smith | smith@math | 19 | 3.8 |
| 53831 | Madayan | madayan@music | 11 | 1.8 |
| 53832 | Guldu | guldu@music | 12 | 2.0 |



An Example Table

- Student is a **relation** on $\text{Int} \times \text{String} \times \text{String} \times \text{Int} \times \text{Float}$
- **Row** represents a tuple (related values)



| <u>sid</u> | name | login | age | gpa |
|------------|---------|---------------|-----|-----|
| 50000 | Dave | dave@cs | 19 | 3.3 |
| 53666 | Jones | jones@cs | 18 | 3.4 |
| 53688 | Smith | smith@ee | 18 | 3.2 |
| 53650 | Smith | smith@math | 19 | 3.8 |
| 53831 | Madayan | madayan@music | 11 | 1.8 |
| 53832 | Guldu | guldu@music | 12 | 2.0 |



Another example: Courses

- Course (cid, instructor, semester)

| <u>cid</u> | instructor | semester |
|------------|------------|----------|
| Piano101 | Jane | Fall 12 |
| Jazz203 | Bob | Sum 12 |
| Calc101 | Mary | Spr 12 |
| Hist105 | Alice | Fall 12 |



Keys

- **Primary key** – minimal subset of fields that *uniquely identifies a tuple*
 - sid is primary key for Students
 - cid is primary key for Courses

primary key

| <u>sid</u> | name | login | age | gpa |
|------------|-------|----------|-----|-----|
| 50000 | Dave | dave@cs | 19 | 3.3 |
| 53666 | Jones | jones@cs | 18 | 3.4 |
| 53688 | Smith | smith@ee | 18 | 3.2 |



Keys

- If we know that login values are unique for all students, login could also be used as a key, which is called a **candidate key**

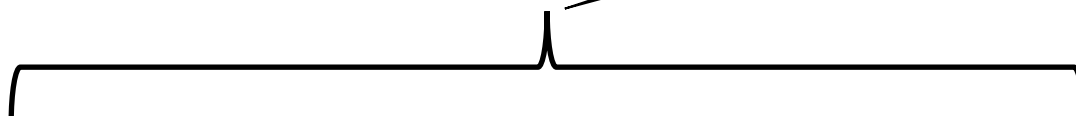
Candidate key

| <u>sid</u> | name | login | age | gpa |
|------------|-------|----------|-----|-----|
| 50000 | Dave | dave@cs | 19 | 3.3 |
| 53666 | Jones | jones@cs | 18 | 3.4 |
| 53688 | Smith | smith@ee | 18 | 3.2 |

Keys

- A (minimum) set of attributes that uniquely identifies tuples may also be used as a candidate key; it is called a **composite key**

Composite key



| fid | date | seat | fname | lname |
|-----|---------|------|-------|----------|
| 734 | 4/5/11 | 22A | Joe | Smith |
| 734 | 5/19/11 | 22A | Peggy | Brooks |
| 734 | 5/20/11 | 22A | Mary | Holcombe |



Table relationships

- Tables can be related, representing dependencies among tuples
- For example, a Course is offered by a Department
- A tuple (row) in one table must identify (reference) a related tuple in another table
- This is done with the use of a **foreign key**
- A foreign key references a primary key in the other table

Table relationships

Foreign key – used for relationships between tables

- Course (cid, instructor, quarter, deptid)
- Department (did, name, college)

extra attribute

primary key

foreign key

Course

| <u>cid</u> | instr | sem | deptid |
|------------|-------|--------|--------|
| Piano101 | Jane | Fall12 | 101 |
| Jazz203 | Bob | Sum12 | 101 |
| Calc101 | Mary | Spr12 | 102 |
| Hist105 | Alice | Fall12 | 103 |

offeredBy

Department

| <u>did</u> | name | college |
|------------|---------|------------|
| 101 | Music | Arts & Sci |
| 102 | Math | Arts & Sci |
| 103 | History | Arts & Sci |
| 112 | Econ | Business |

Many to many relationships

- In general, we need a new table **Enrolls(crsid, studid)**
crsid is a foreign key that references cid in the Course table
studid is a foreign key that references sid in the Student table

Course

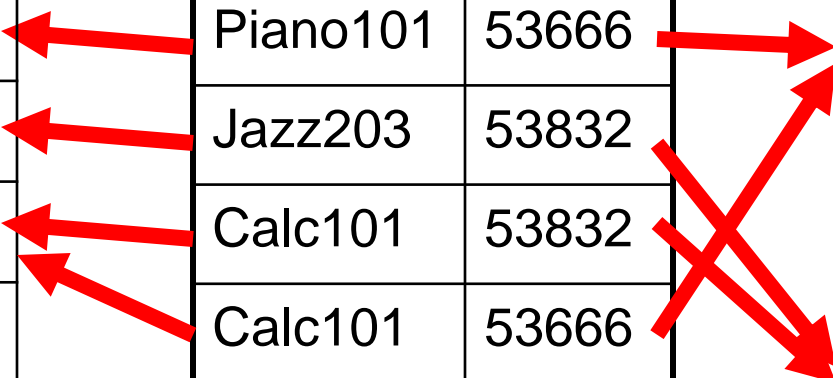
| <u>cid</u> | instr | sem |
|------------|-------|--------|
| Piano101 | Jane | Fall19 |
| Jazz203 | Bob | Sum19 |
| Calc101 | Mary | Spr19 |
| Hist105 | Alice | Fall19 |

Enrolls

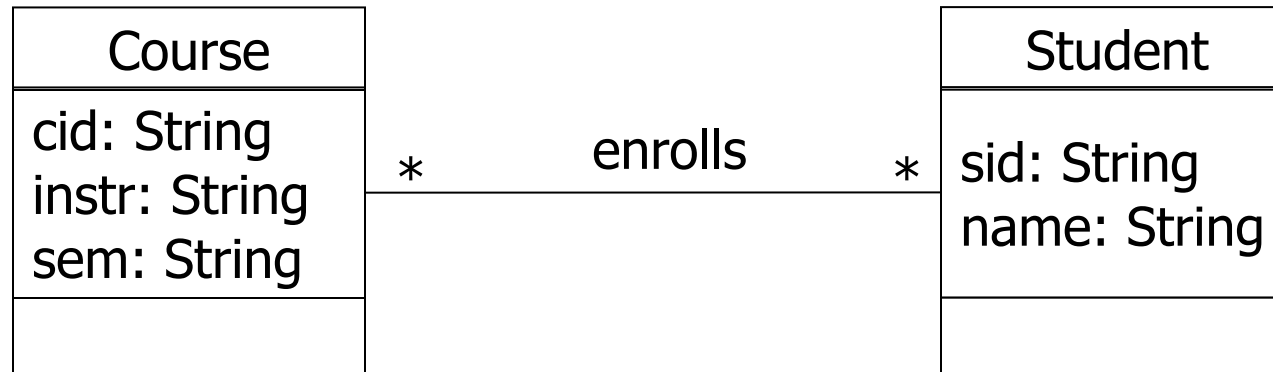
| crsid | studid |
|----------|--------|
| Piano101 | 53666 |
| Jazz203 | 53832 |
| Calc101 | 53832 |
| Calc101 | 53666 |

Student

| <u>sid</u> | name |
|------------|-------|
| 50000 | Dave |
| 53666 | Jones |
| 53688 | Smith |
| 53650 | Smith |
| 53831 | Patel |
| 53832 | Zak |



Relational tables and class diagrams



| <u>cid</u> | instr | sem |
|------------|-------|--------|
| Piano101 | Jane | Fall19 |
| Jazz203 | Bob | Sum19 |
| Calc101 | Mary | Spr19 |
| Hist105 | Alice | Fall19 |

| crsid | studid |
|----------|--------|
| Piano101 | 53666 |
| Jazz203 | 53832 |
| Calc101 | 53832 |
| Calc101 | 53666 |

| <u>sid</u> | name |
|------------|-------|
| 50000 | Dave |
| 53666 | Jones |
| 53688 | Smith |
| 53650 | Smith |
| 53831 | Patel |
| 53832 | Zak |

Red arrows indicate the mapping from the Enrolls table to the Course and Student tables. Arrows point from the crsid column to the cid column and from the studid column to the sid column. A red 'X' is drawn over the bottom two rows of the Enrolls table, indicating a conflict or error in the mapping.



Relational Algebra

- Created by Codd
- Collection of operators for specifying queries
- Query describes step-by-step procedure for computing answer (i.e., **operational**)
- Each operator accepts one or two relations as input and returns a relation as output
- Relational algebra expression composed of multiple operators



Basic operators

- Selection – return *rows* that meet some condition
- Projection – return *column* values
- Union
- Cross product
- Difference
- Other operators can be defined in terms of basic operators

We will only *outline* a few of them



Example Schema (simplified)

- Course (cid, instructor, quarter, dept)
- Student (sid, name, gpa)
- Enrolls (cid, grade, studid)



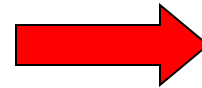
Selection

Select students with gpa higher than 3.3 from S1:

$$\sigma_{gpa > 3.3}(S1)$$

S1

| sid | name | gpa |
|-------|---------|-----|
| 50000 | Dave | 3.3 |
| 53666 | Jones | 3.4 |
| 53688 | Smith | 3.2 |
| 53650 | Smith | 3.8 |
| 53831 | Madayan | 1.8 |
| 53832 | Guldu | 2.0 |



| sid | name | gpa |
|-------|-------|-----|
| 53666 | Jones | 3.4 |
| 53650 | Smith | 3.8 |



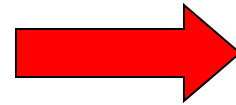
Projection

Project name and gpa of all students in S1:

$$\Pi_{\text{name, gpa}}(S1)$$

S1

| Sid | name | gpa |
|-------|---------|-----|
| 50000 | Dave | 3.3 |
| 53666 | Jones | 3.4 |
| 53688 | Smith | 3.2 |
| 53650 | Smith | 3.8 |
| 53831 | Madayan | 1.8 |
| 53832 | Guldu | 2.0 |



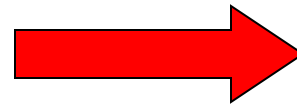
| name | gpa |
|---------|-----|
| Dave | 3.3 |
| Jones | 3.4 |
| Smith | 3.2 |
| Smith | 3.8 |
| Madayan | 1.8 |
| Guldu | 2.0 |

Combine Selection and Projection

Project name and gpa of students in S1 with gpa > 3.3:

$$\Pi_{\text{name,gpa}} (\sigma_{\text{gpa} > 3.3} (S1))$$

| Sid | name | gpa |
|-------|---------|-----|
| 50000 | Dave | 3.3 |
| 53666 | Jones | 3.4 |
| 53688 | Smith | 3.2 |
| 53650 | Smith | 3.8 |
| 53831 | Madayan | 1.8 |
| 53832 | Guldu | 2.0 |



| name | gpa |
|-------|-----|
| Jones | 3.4 |
| Smith | 3.8 |



Joins

- Combine information from two or more tables
- Example: courses owned by departments:

$C1 \bowtie_{C1.deptid = D.did} D$

C1

| <u>cid</u> | instr | sem | deptid |
|------------|-------|--------|--------|
| Piano101 | Jane | Fall12 | 101 |
| Jazz203 | Bob | Sum12 | 101 |
| Calc101 | Mary | Spr12 | 102 |
| Hist105 | Alice | Fall12 | 103 |

D

| <u>did</u> | name | college |
|------------|---------|------------|
| 101 | Music | Arts & Sci |
| 102 | Math | Arts & Sci |
| 103 | History | Arts & Sci |
| 112 | Econ | Business |

Joins

C1

| <u>cid</u> | instr | sem | deptid |
|------------|-------|--------|--------|
| Piano101 | Jane | Fall12 | 101 |
| Jazz203 | Bob | Sum12 | 101 |
| Calc101 | Mary | Spr12 | 102 |
| Hist105 | Alice | Fall12 | 103 |

D

| <u>did</u> | name | college |
|------------|---------|------------|
| 101 | Music | Arts & Sci |
| 102 | Math | Arts & Sci |
| 103 | History | Arts & Sci |
| 112 | Econ | Business |

| <u>cid</u> | instr | sem | deptid | <u>did</u> | name | college |
|------------|-------|--------|--------|------------|---------|------------|
| Piano101 | Jane | Fall12 | 101 | 101 | Music | Arts & Sci |
| Jazz203 | Bob | Sum12 | 101 | 101 | Music | Arts & Sci |
| Calc101 | Mary | Spr12 | 102 | 102 | Math | Arts & Sci |
| Hist105 | Alice | Fall12 | 103 | 103 | History | Arts & Sci |



History of SQL

- In 1974, D. Chamberlin (IBM San Jose Laboratory) defined language called 'Structured English Query Language' (SEQUEL).
- A revised version, SEQUEL/2, was defined in 1976 but name was subsequently changed to SQL for legal reasons.
- Still pronounced 'see-quel', though official pronunciation is 'S-Q-L'.
- IBM subsequently produced a prototype DBMS called *System R*, based on SEQUEL/2.
- Roots of SQL, however, are in SQUARE (Specifying Queries as Relational Expressions), which predates System R project.



History of SQL

- In late 70s, ORACLE was introduced as (likely) the first commercial RDBMS based on SQL
- In 1987, ANSI and ISO published an initial standard for SQL
- In 1992, first major revision to ISO standard occurred, referred to as SQL2 or SQL/92
- In 1999, SQL3 was released with support for object-oriented data management



Objectives of SQL

- SQL includes two major components:
 - A **Data Definition Language (DDL)** for defining a database structure
Create table, create indexes, alter table, etc.
 - A **Data Manipulation Language (DML)** for retrieving and updating data
Select rows from one or more tables (using joins), insert, update, or delete rows



Intro to SQL

- CREATE TABLE
 - Create a new table, e.g., students, courses, enrolled
- Also, ALTER TABLE, DROP TABLE, and other statements
- CREATE INDEX and DROP INDEX
- SELECT-FROM-WHERE
 - List all CS courses
- INSERT
 - Add new students, courses, or enroll students in courses
- UPDATE
 - Update attributes of students, courses, or change student enrollments
- DELETE
 - Delete students or courses



Create Table

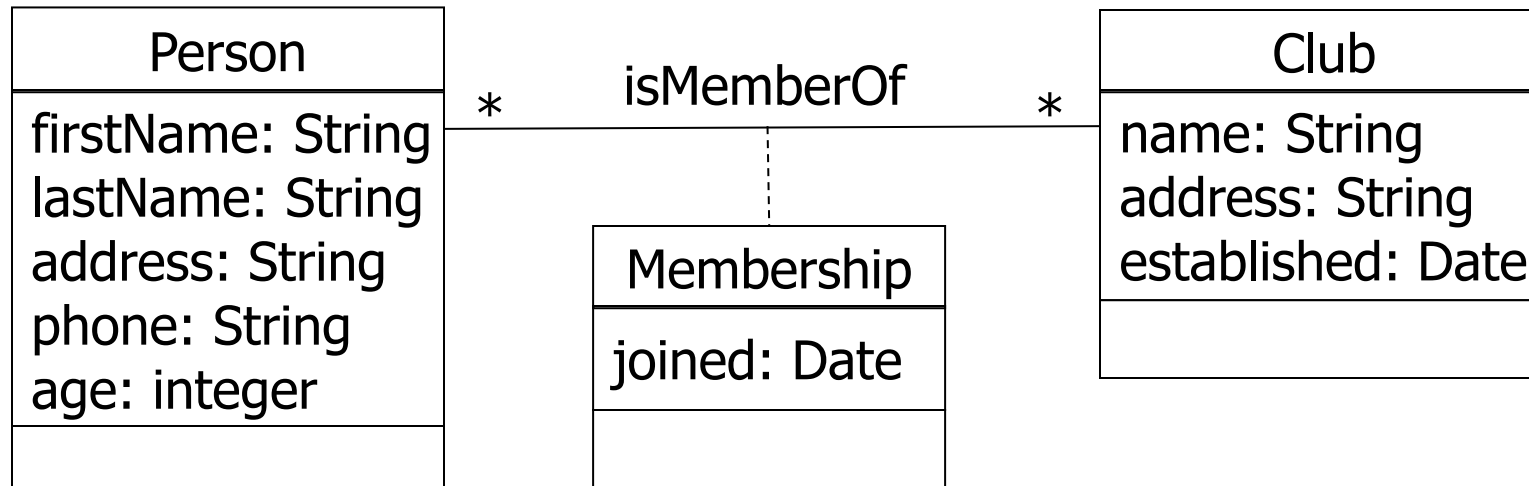
```
CREATE TABLE TableName
    {(colName dataType [NOT NULL] [UNIQUE]
        [DEFAULT defaultOption]
        [CHECK searchCondition] [,...]}
        [PRIMARY KEY (listOfColumns),]
    {[UNIQUE (listOfColumns),] [...,]}
    {[FOREIGN KEY (listOfFKColumns)
        REFERENCES ParentTableName [(listOfCKColumns)],
        [ON UPDATE referentialAction]
        [ON DELETE referentialAction ]] [,...]}
    {[CHECK (searchCondition)] [,...]} }
```



Create Table

- Creates a table with one or more columns of the specified *dataType*.
- With NOT NULL, system rejects any attempt to insert a NULL value in the column.
- Can specify a DEFAULT value for the column.
- Primary keys should always be specified as NOT NULL.
- FOREIGN KEY clause specifies FK along with the referential action

Create Table



```
CREATE TABLE Person (  
    Id                INT                UNSIGNED PRIMARY KEY,  
    FirstName         VARCHAR(255)      NOT NULL,  
    LastName          VARCHAR(255)      NOT NULL,  
    Address            VARCHAR(255),  
    Phone              VARCHAR(255),  
    Age               INT UNSIGNED  
);
```




Create Table

```
CREATE TABLE Club (  
    Id            INT            UNSIGNED PRIMARY KEY,  
    Name          VARCHAR(255)   NOT NULL,  
    Address        VARCHAR(255),  
    Established    DATETIME  
);
```

```
CREATE TABLE IsMemberOf (  
    Id            INT            UNSIGNED PRIMARY KEY,  
    PersonId      INT UNSIGNED   NOT NULL,  
    ClubId        INT UNSIGNED   NOT NULL,  
    Joined        DATETIME,  
  
    FOREIGN KEY   (PersonId) REFERENCES Person(Id),  
    FOREIGN KEY   (ClubId)   REFERENCES Club(Id)  
);
```



SELECT Statement

```
SELECT [DISTINCT | ALL]
      { * | [columnExpression [AS newName]] [, ...] }
FROM  TableName [alias] [, ...]
[WHERE condition]
[GROUP BY columnList] [HAVING condition]
[ORDER BY columnList]
```

- Order of the clauses cannot be changed
- Only SELECT and FROM are mandatory



SELECT Statement

| | |
|----------|--|
| FROM | Specifies table(s) to be used. |
| WHERE | Filters rows. |
| GROUP BY | Forms groups of rows with same column value. |
| HAVING | Filters groups subject to some condition. |
| SELECT | Specifies which columns are to appear in output. |
| ORDER BY | Specifies the order of the output. |



Select-From-Where query

“Find all clubs”

```
SELECT * FROM Club
```

The * above means “get all column values”

“Find all persons who are under 18”

```
SELECT *  
FROM   Person p  
WHERE  p.Age < 18
```



Select-From-Where query

“Find names of all persons who are under 18”

```
SELECT  p.FirstName  
FROM    Person p  
WHERE   p.Age < 18
```

The above query performs a **selection and projection**

Queries across multiple tables (joins)

“Print names and address of persons younger than 20 who are members of the Tennis club”

Person

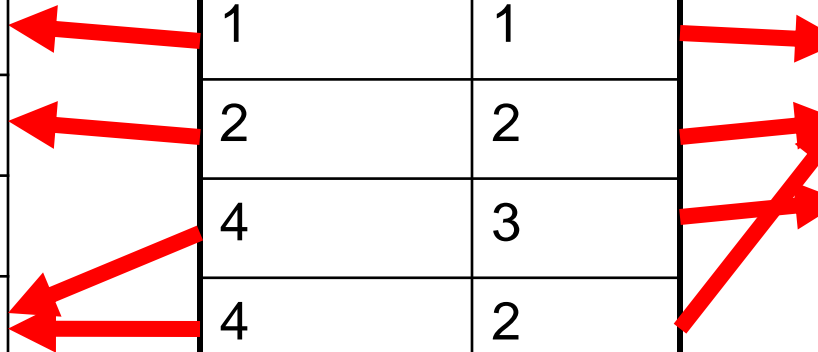
| <u>Id</u> | FirstName | Age |
|-----------|-----------|-----|
| 1 | Jane | 12 |
| 2 | Bob | 22 |
| 3 | Mary | 31 |
| 4 | Alice | 17 |

IsMemberOf

| <u>PersonId</u> | ClubId |
|-----------------|--------|
| 1 | 1 |
| 2 | 2 |
| 4 | 3 |
| 4 | 2 |

Club

| <u>Id</u> | Name |
|-----------|--------|
| 1 | Chess |
| 2 | Tennis |
| 3 | Bridge |
| 4 | Swim |





Queries across multiple tables (joins)

“Print names and address of persons younger than 20 who are members of the Tennis club”

```
SELECT  p.FirstName, p.Address
FROM    Person p, Club c, IsMemberOf m
WHERE   p.Age < 20  AND  c.Name = 'Tennis'
        AND  m.PersonId = p.Id
        AND  m.ClubId = c.Id
```



Queries across multiple tables (joins)

“Print names and address of persons younger than 20 who are members of the Tennis club”

```
SELECT  p.FirstName, p.Address
FROM    Person p, Club c, IsMemberOf m
WHERE   p.Age < 20 AND c.Name = 'Tennis'
        AND m.PersonId = p.Id
        AND m.ClubId = c.Id
```

an alias

join clauses



INSERT

```
INSERT INTO TableName [ (columnList) ]  
VALUES (dataValueList)
```

- **columnList** is optional; if omitted, SQL assumes a list of all columns in their original CREATE TABLE order
- any columns omitted must have been declared as NULL when table was created, unless DEFAULT was specified when creating column



INSERT

- **dataValueList** must match **columnList** as follows:
 - number of items in each list must be same;
 - must be direct correspondence in position of items in two lists;
 - data type of each item in **dataValueList** must be compatible with data type of corresponding column

```
INSERT INTO Person (FirstName, LastName, Address,  
Phone, Age) VALUES ( 'Jeff', 'Roberts', '11 Oak St',  
'123-444-5566', 24 )
```



UPDATE

```
UPDATE TableName  
SET  columnName1 = dataValue1  
    [, columnName2 = dataValue2...]  
[WHERE searchCondition]
```

- **TableName** can be name of a base table or an updatable view
- SET clause specifies names of one or more columns that are to be updated



UPDATE

- WHERE clause is optional:
 - if omitted, named columns are updated for all rows in table;
 - if specified, only those rows that satisfy **searchCondition** are updated.
- New **dataValue**(s) must be compatible with data type for corresponding column



DELETE

DELETE FROM TableName
[WHERE searchCondition]

- **TableName** can be name of a base table or an updatable view.
- **searchCondition** is optional; if omitted, all rows are deleted from table. This does not delete table. If **searchCondition** is specified, only those rows that satisfy condition are deleted.



Examples of other SQL statements

```
INSERT INTO Club (Name, Address, Established)
VALUES ( 'Chess', '33 Leaf St., Blossom, OR. 88888',
        '2007-07-12 12:00:00' )
```

```
UPDATE Club
SET Address = '11 Trunk St., Blossom, OR. 77777'
WHERE Name = 'Chess'
```

```
DELETE FROM Club
WHERE Name = 'Chess'
```



Other SQL features

- MIN, MAX, AVG
 - Find highest grade in fall database course
- COUNT, DISTINCT
 - How many students enrolled in CS courses in the fall?
- ORDER BY, GROUP BY
 - Rank students by their grade in fall database course
- Transactions



Relational DBMS

Examples of commercial RDBMS systems:

ORACLE

DB2 (IBM)

SQL Server (Microsoft)

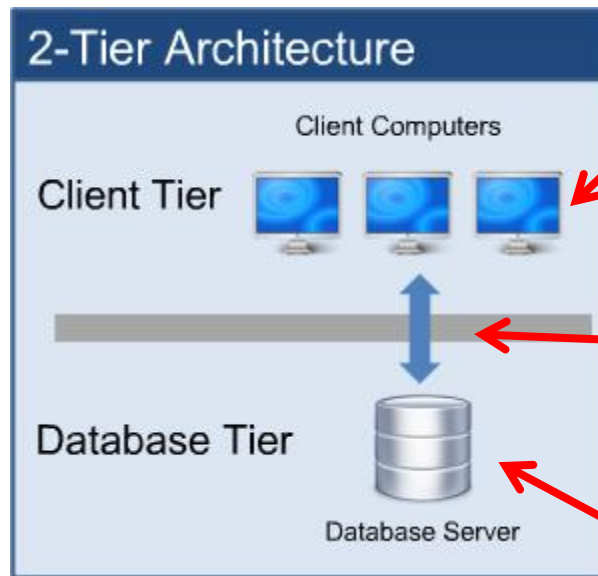
Examples of open source RDBMS systems:

MySQL

PostgreSQL

MSQL

Relational DBMS

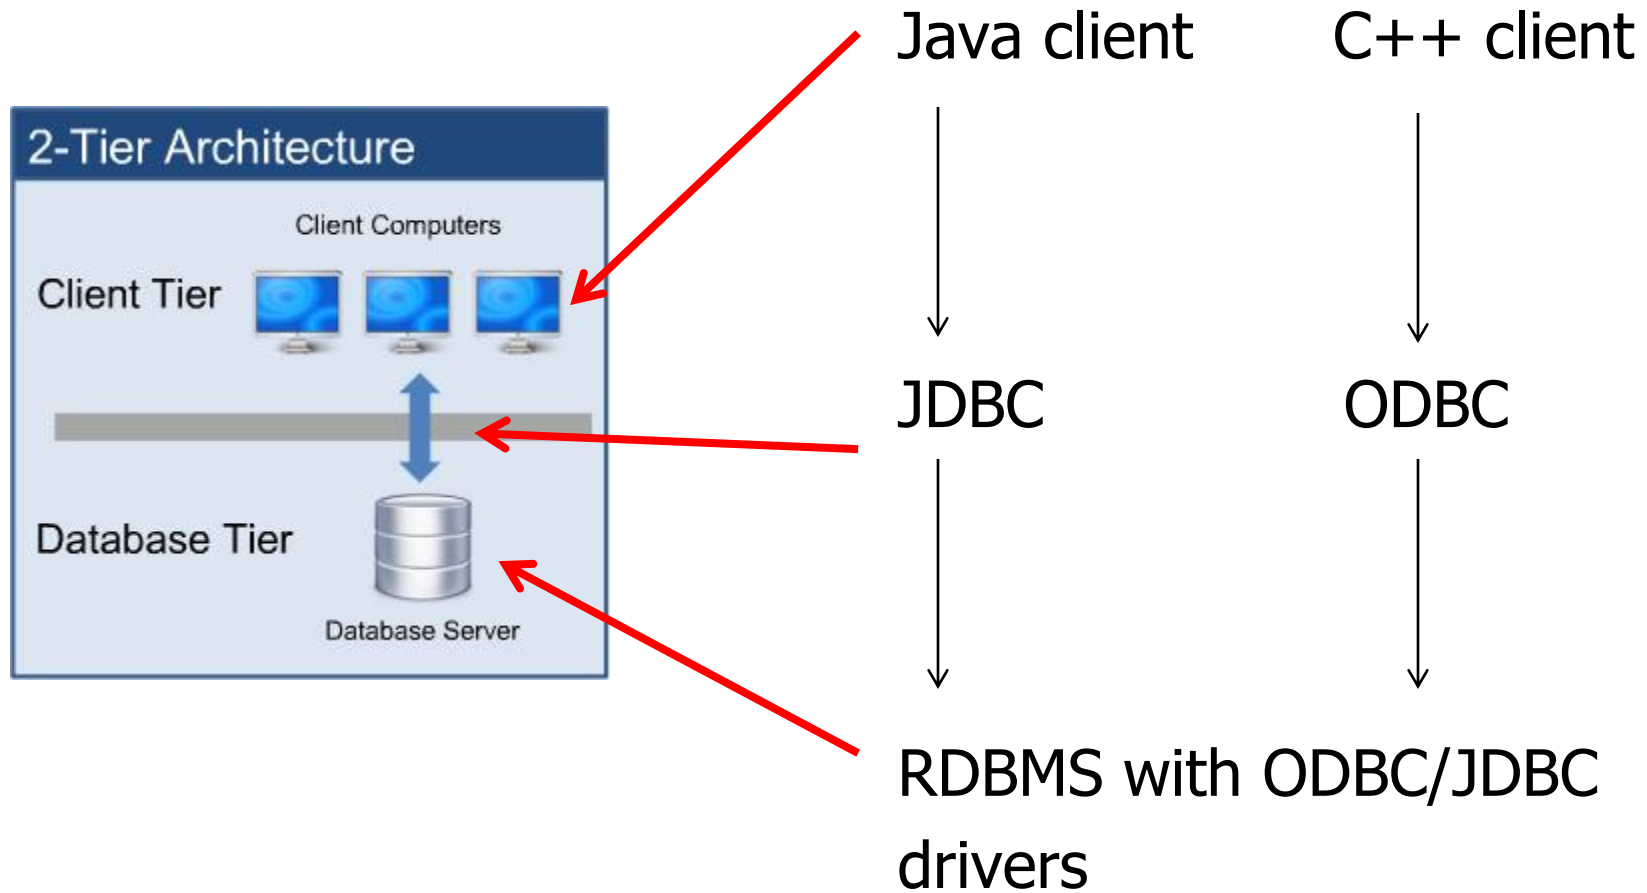


Client program can be in any language, but must "speak" the RDBMS communication language, usually over the network

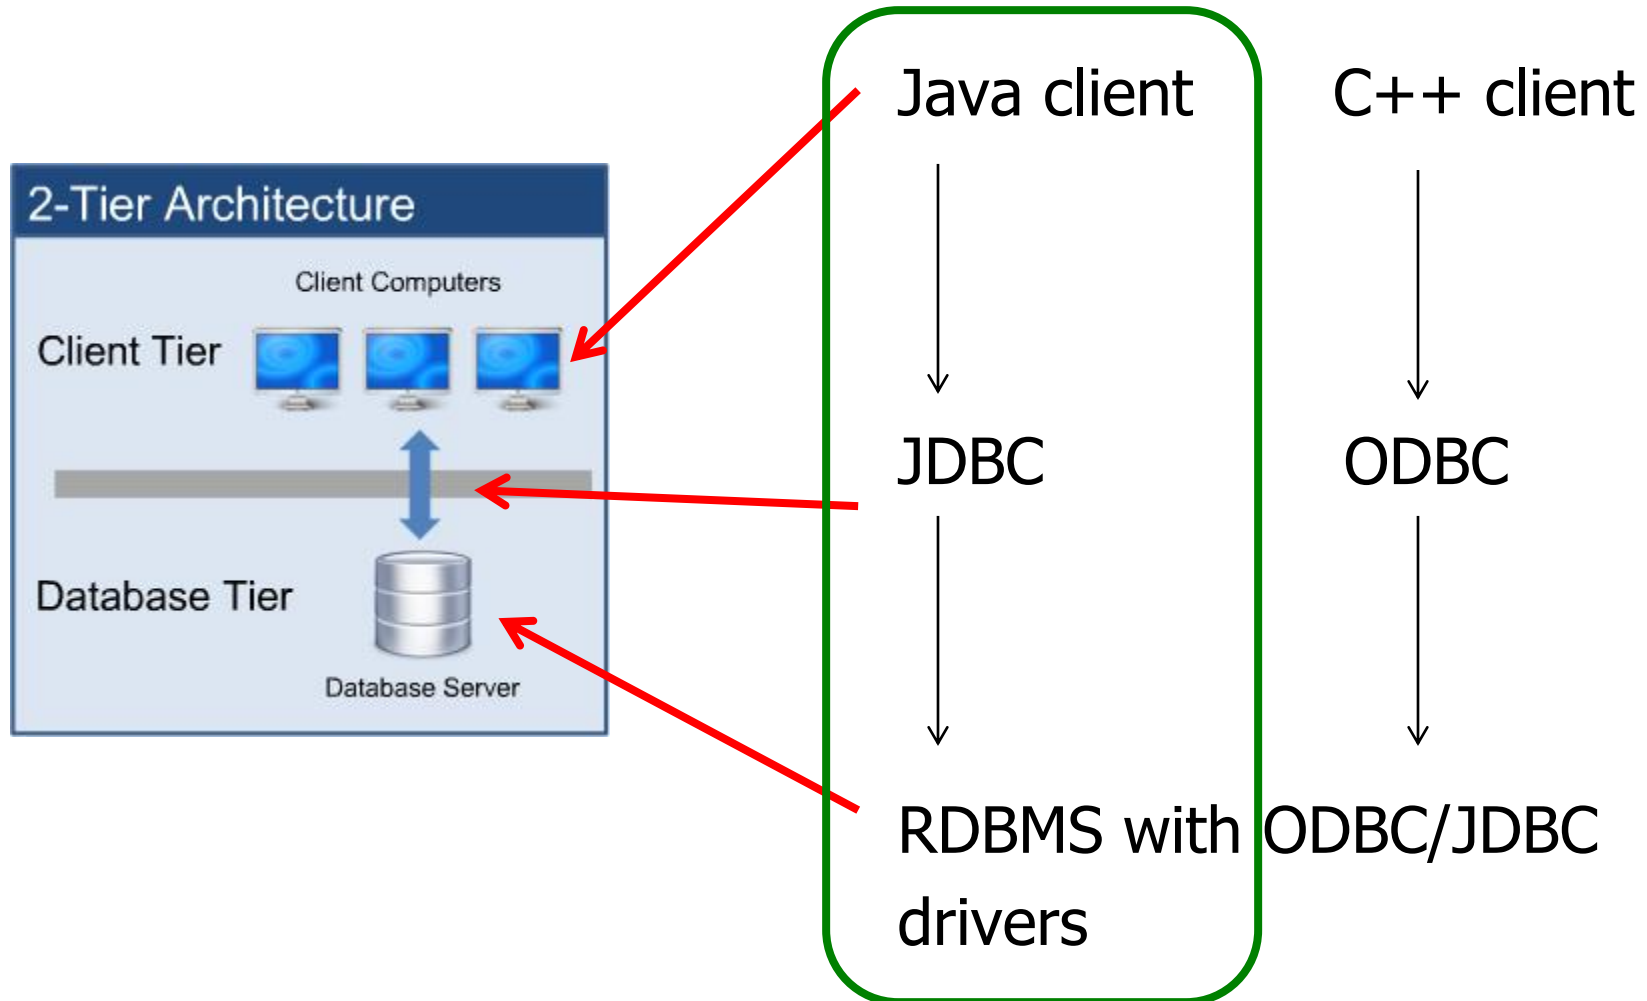
Communication language/protocol or an API

RDBMS server, which accepts request over the network

Relational DBMS



Relational DBMS





JDBC

- JDBC is a Sun (now Oracle) trademark
 - It is often taken to stand for Java Database Connectivity
- Java is standardized, but there are many versions of RDBMS servers
- JDBC is a standard way of accessing SQL databases from Java



Driver types

- There are four types of drivers:
 - **JDBC Type 1 Driver** -- JDBC/ODBC Bridge drivers
 - ODBC (Open DataBase Connectivity) is a standard software API designed to be independent of specific programming languages
 - Sun provides a JDBC/ODBC implementation
 - **JDBC Type 2 Driver** -- use platform-specific APIs for data access
 - **JDBC Type 3 Driver** -- 100% Java, use a net protocol to access a remote listener and map calls into vendor-specific calls
 - **JDBC Type 4 Driver** -- 100% Java



Connector/J

- We will be using MySQL open-source RDBMS
- MySQL documentation:
<https://dev.mysql.com/doc/refman/5.1/en/index.html>
- Connector/J is a JDBC Type 4 Driver for connecting Java to MySQL
- It is available as a jar file in /opt/classes
`mysql-connector-java-5.1.26-bin.jar`
- As usual, it must be available on the Java CLASSPATH



Typical use of JDBC

1. Establish a **connection**
2. Create a JDBC **Statement**
3. **Execute** an SQL Statement
4. GET the **ResultSet**
5. Process the results from the ResultSet
6. **Close** the connection



Mapping types JDBC - Java

Impedance mismatch: differences between an OO language world and an RDB world;
for example, data type differences:

| JDBC Type | Java Type |
|--------------------------------------|-----------|
| BIT | boolean |
| TINYINT | byte |
| SMALLINT | short |
| INTEGER | int |
| BIGINT | long |
| REAL | float |
| FLOAT | double |
| DOUBLE | |
| BINARY VARBINARY LONGVARBINARY | byte[] |
| CHAR VARCHAR LONGVARCHAR | String |

| JDBC Type | Java Type |
|--------------------|----------------------------|
| NUMERIC DECIMAL | BigDecimal |
| DATE | java.sql.Date |
| TIME TIMESTAMP | java.sql.Timestamp |
| CLOB | Clob* |
| BLOB | Blob* |
| ARRAY | Array* |
| DISTINCT | mapping of underlying type |
| STRUCT | Struct* |
| REF | Ref* |
| JAVA_OBJECT | underlying Java class |

*SQL3 data type supported in JDBC 2.0



Object persistence

- Objects may be written to and retrieved from disk using serialization; a class needs to implement the **serializable** interface (even with no methods)

```
public class Data implements serializable  
{ ... }
```

- **ObjectOutputStream** and **ObjectInputStream** should be used with the **writeObject** and **readObject** methods
 - What are the problems with this approach?



Object persistence

- Objects may be written to and retrieved from disk using serialization; a class needs to implement the **serializable** interface (even with no methods)

```
public class Data implements serializable
{ ... }
```

- **ObjectOutputStream** and **ObjectInputStream** should be used with the **writeObject** and **readObject** methods
 - What are the problems with this approach?
 - Persisting thousands of objects may result in creating and managing thousands of files. What about millions of objects?



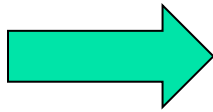
Object persistence

- Objects may be written to and retrieved from a relational database, such as MySQL
- Object-Relational-Mapping (ORM)
 - Classes are mapped onto tables
 - Associations onto foreign keys and/or relation tables
 - A newly created object is inserted into its class's table by storing its state as a single row with its table's attribute values
 - An existing object can be restored by retrieving its state from the database and creating a new instance initialized to the values retrieved from the database

Object persistence

- Classes are mapped onto tables

| Club |
|-------------------|
| name: String |
| address: String |
| established: Date |
| op(arg:int): long |



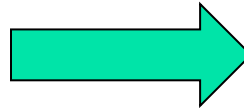
| Club | | varchar(255) | datetime |
|-----------|------|--------------|-------------|
| <u>id</u> | name | address | established |
| ... | ... | ... | ... |
| | | | |
| | | | |

- Preserve the attribute names and select suitable types from SQL
- Define a column to serve as a key (automatically generated)
- Set constraints, if needed (unique, enum-like values, ...)
- Operations are not represented

Object persistence

- Objects are represented as rows

| <u>: Club</u> |
|-------------------------|
| name = "Chess" |
| address = "11 Oak St." |
| established = 4/22/2004 |

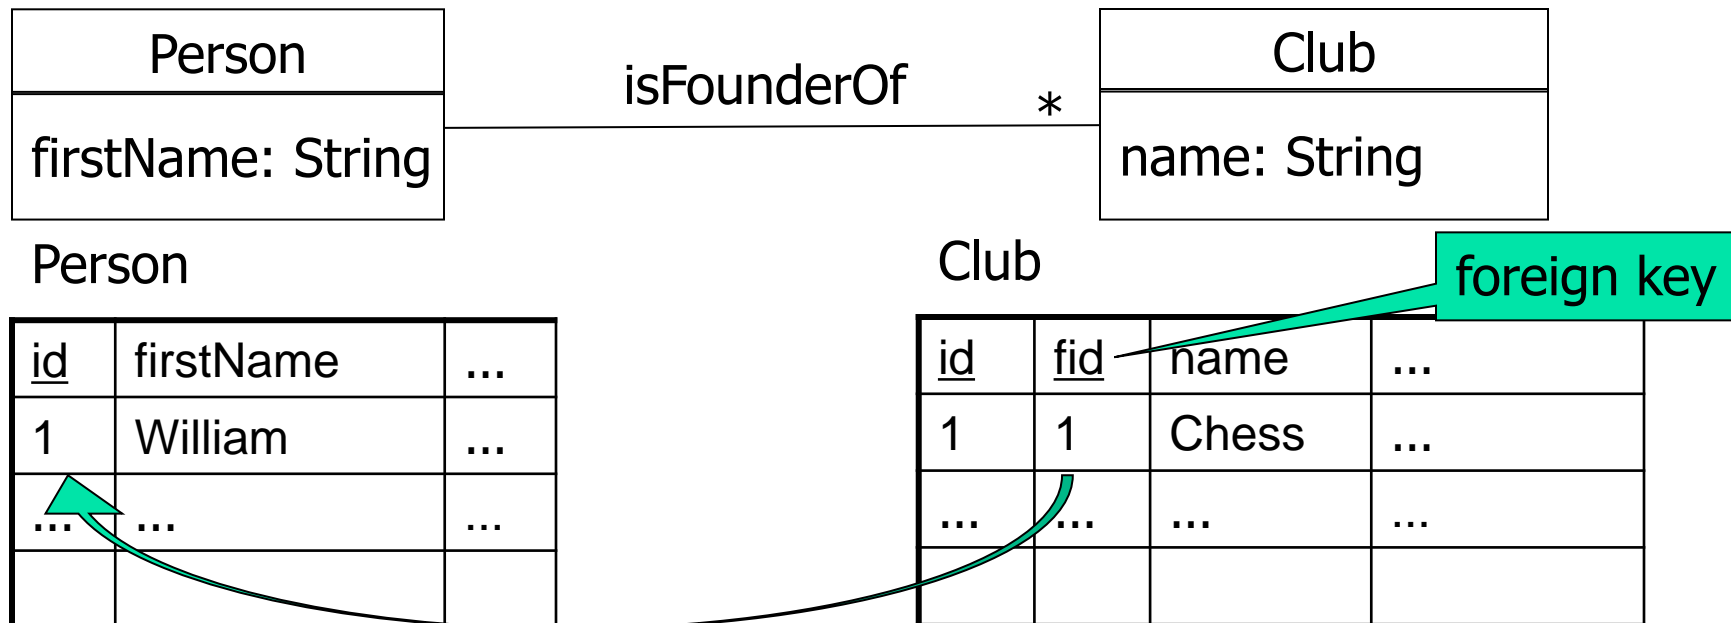


| Club | | varchar(255) | datetime |
|-----------|-------|--------------|-------------|
| <u>id</u> | name | address | established |
| 1 | Chess | 11 Oak St. | 4/22/2004 |
| ... | ... | ... | ... |
| | | | |
| | | | |

- Identifier (primary key) is automatically generated
- Java/C++ data values are automatically converted
- Beware of dates:
java.util.date is not quite the same as java.sql.date !

Object persistence

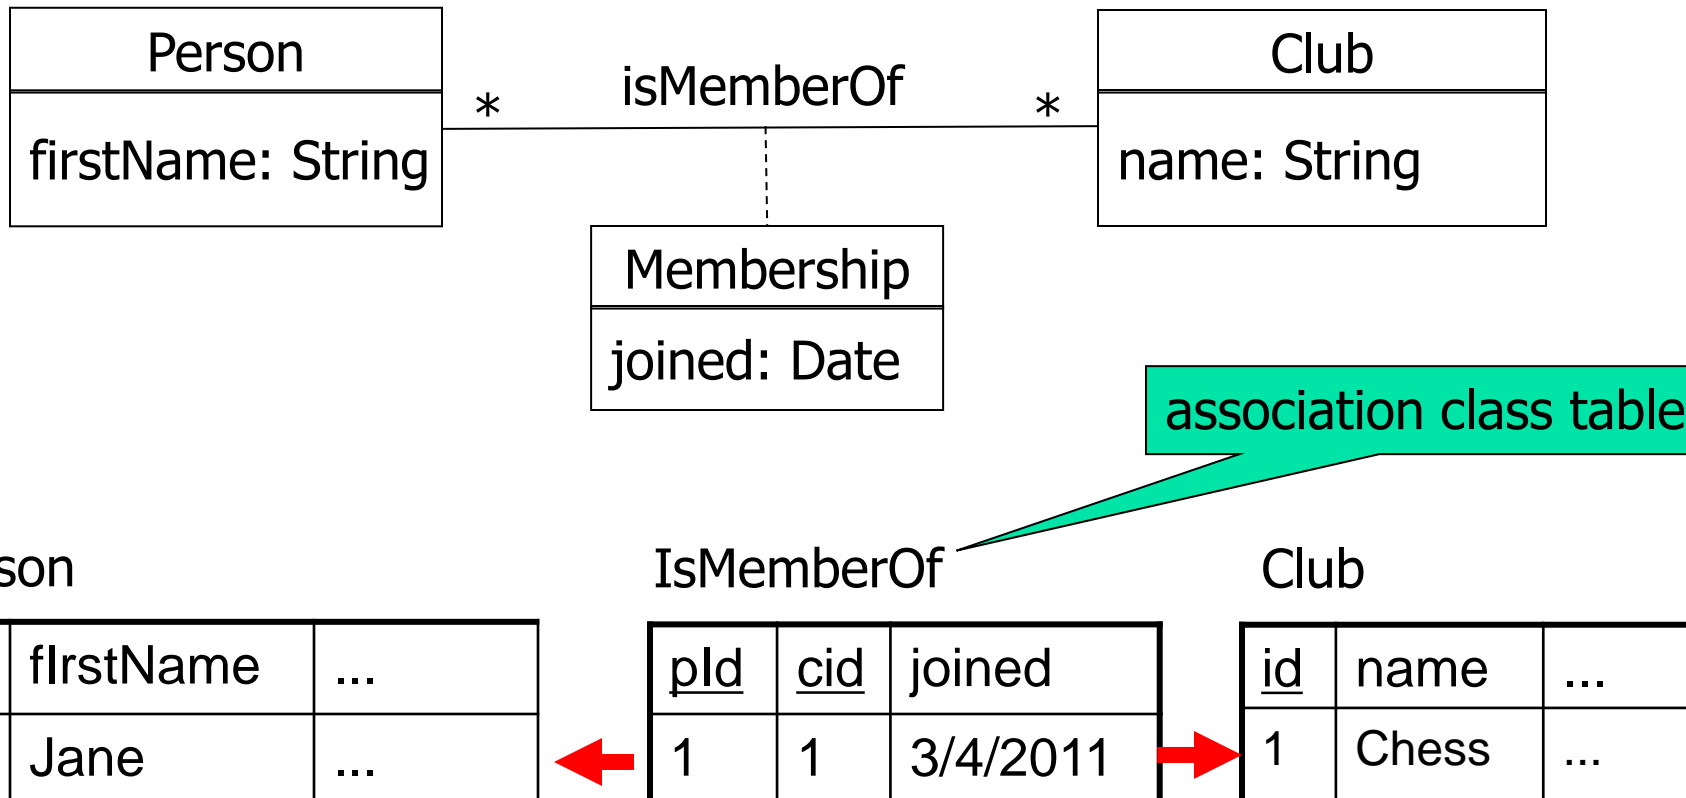
- 1-1 and 1-m associations are mapped onto foreign keys



- Remember, this works for 1-1, 1-optional, and 1-m associations!
- Only one table should have the foreign key defined (for 1-m, it should be the table on the "many" side)

Object persistence

- m-m associations and association classes are mapped onto relation tables



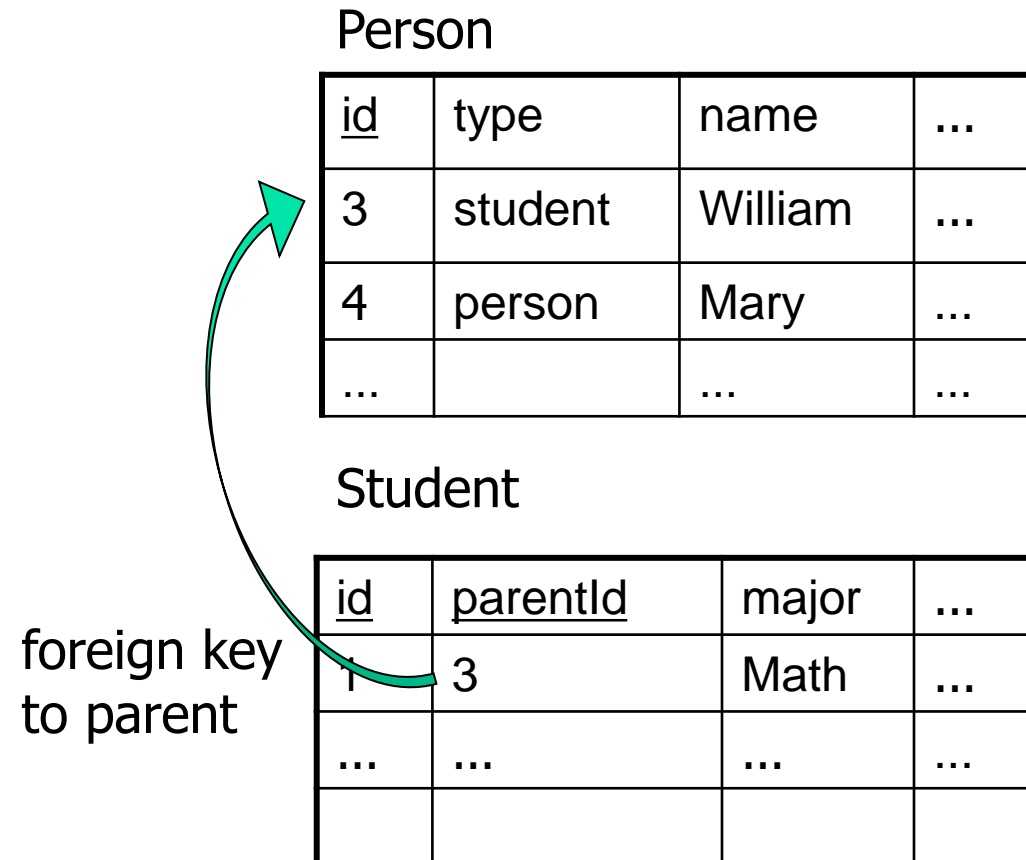
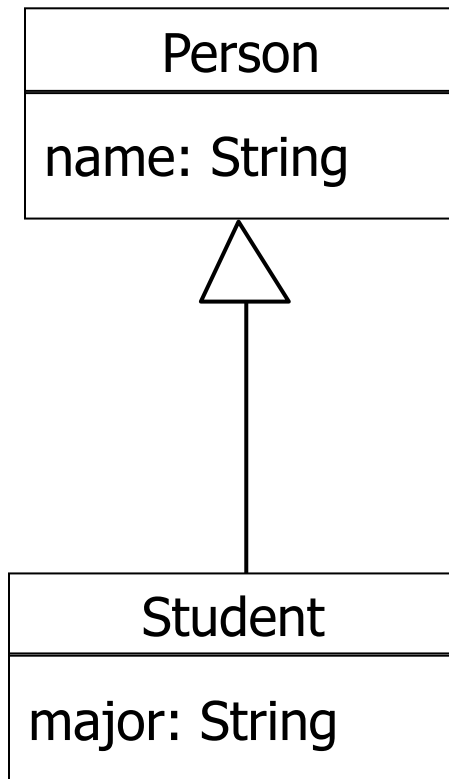


Object persistence

- Mapping of generalization relationships is more involved
- Method 1:
 - Parent and Child classes are mapped onto their own tables
 - The Child class table has a foreign key to the Parent table
 - A Child class object is represented partially in the Child table (Child class attributes) and partially in the Parent class (Parent class attributes)
 - Retrieval of a complete Child class object requires a join SQL statement to retrieve all attributes; the foreign key value connects the two parts
 - Retrieval of objects in a hierarchy requires a **left outer join** of parent and child
 - With large/deep hierarchies this method may be inefficient

Object persistence

- Generalization mapping, method 1



Object persistence

■ Retrieving objects in method 1

select * from
Person INNER JOIN Student on
 Person.id = Student.parentId;
to retrieve just Student objects

OR

select * from
Person LEFT OUTER JOIN Student on
 Person.id = Student.parentId;
to retrieve both Person **and** Student
objects

a 'type tag' used to
distinguish objects
of specific classes

Person

| <u>id</u> | type | name | ... |
|-----------|---------|---------|-----|
| 3 | student | William | ... |
| 4 | person | Mary | ... |
| ... | | ... | ... |

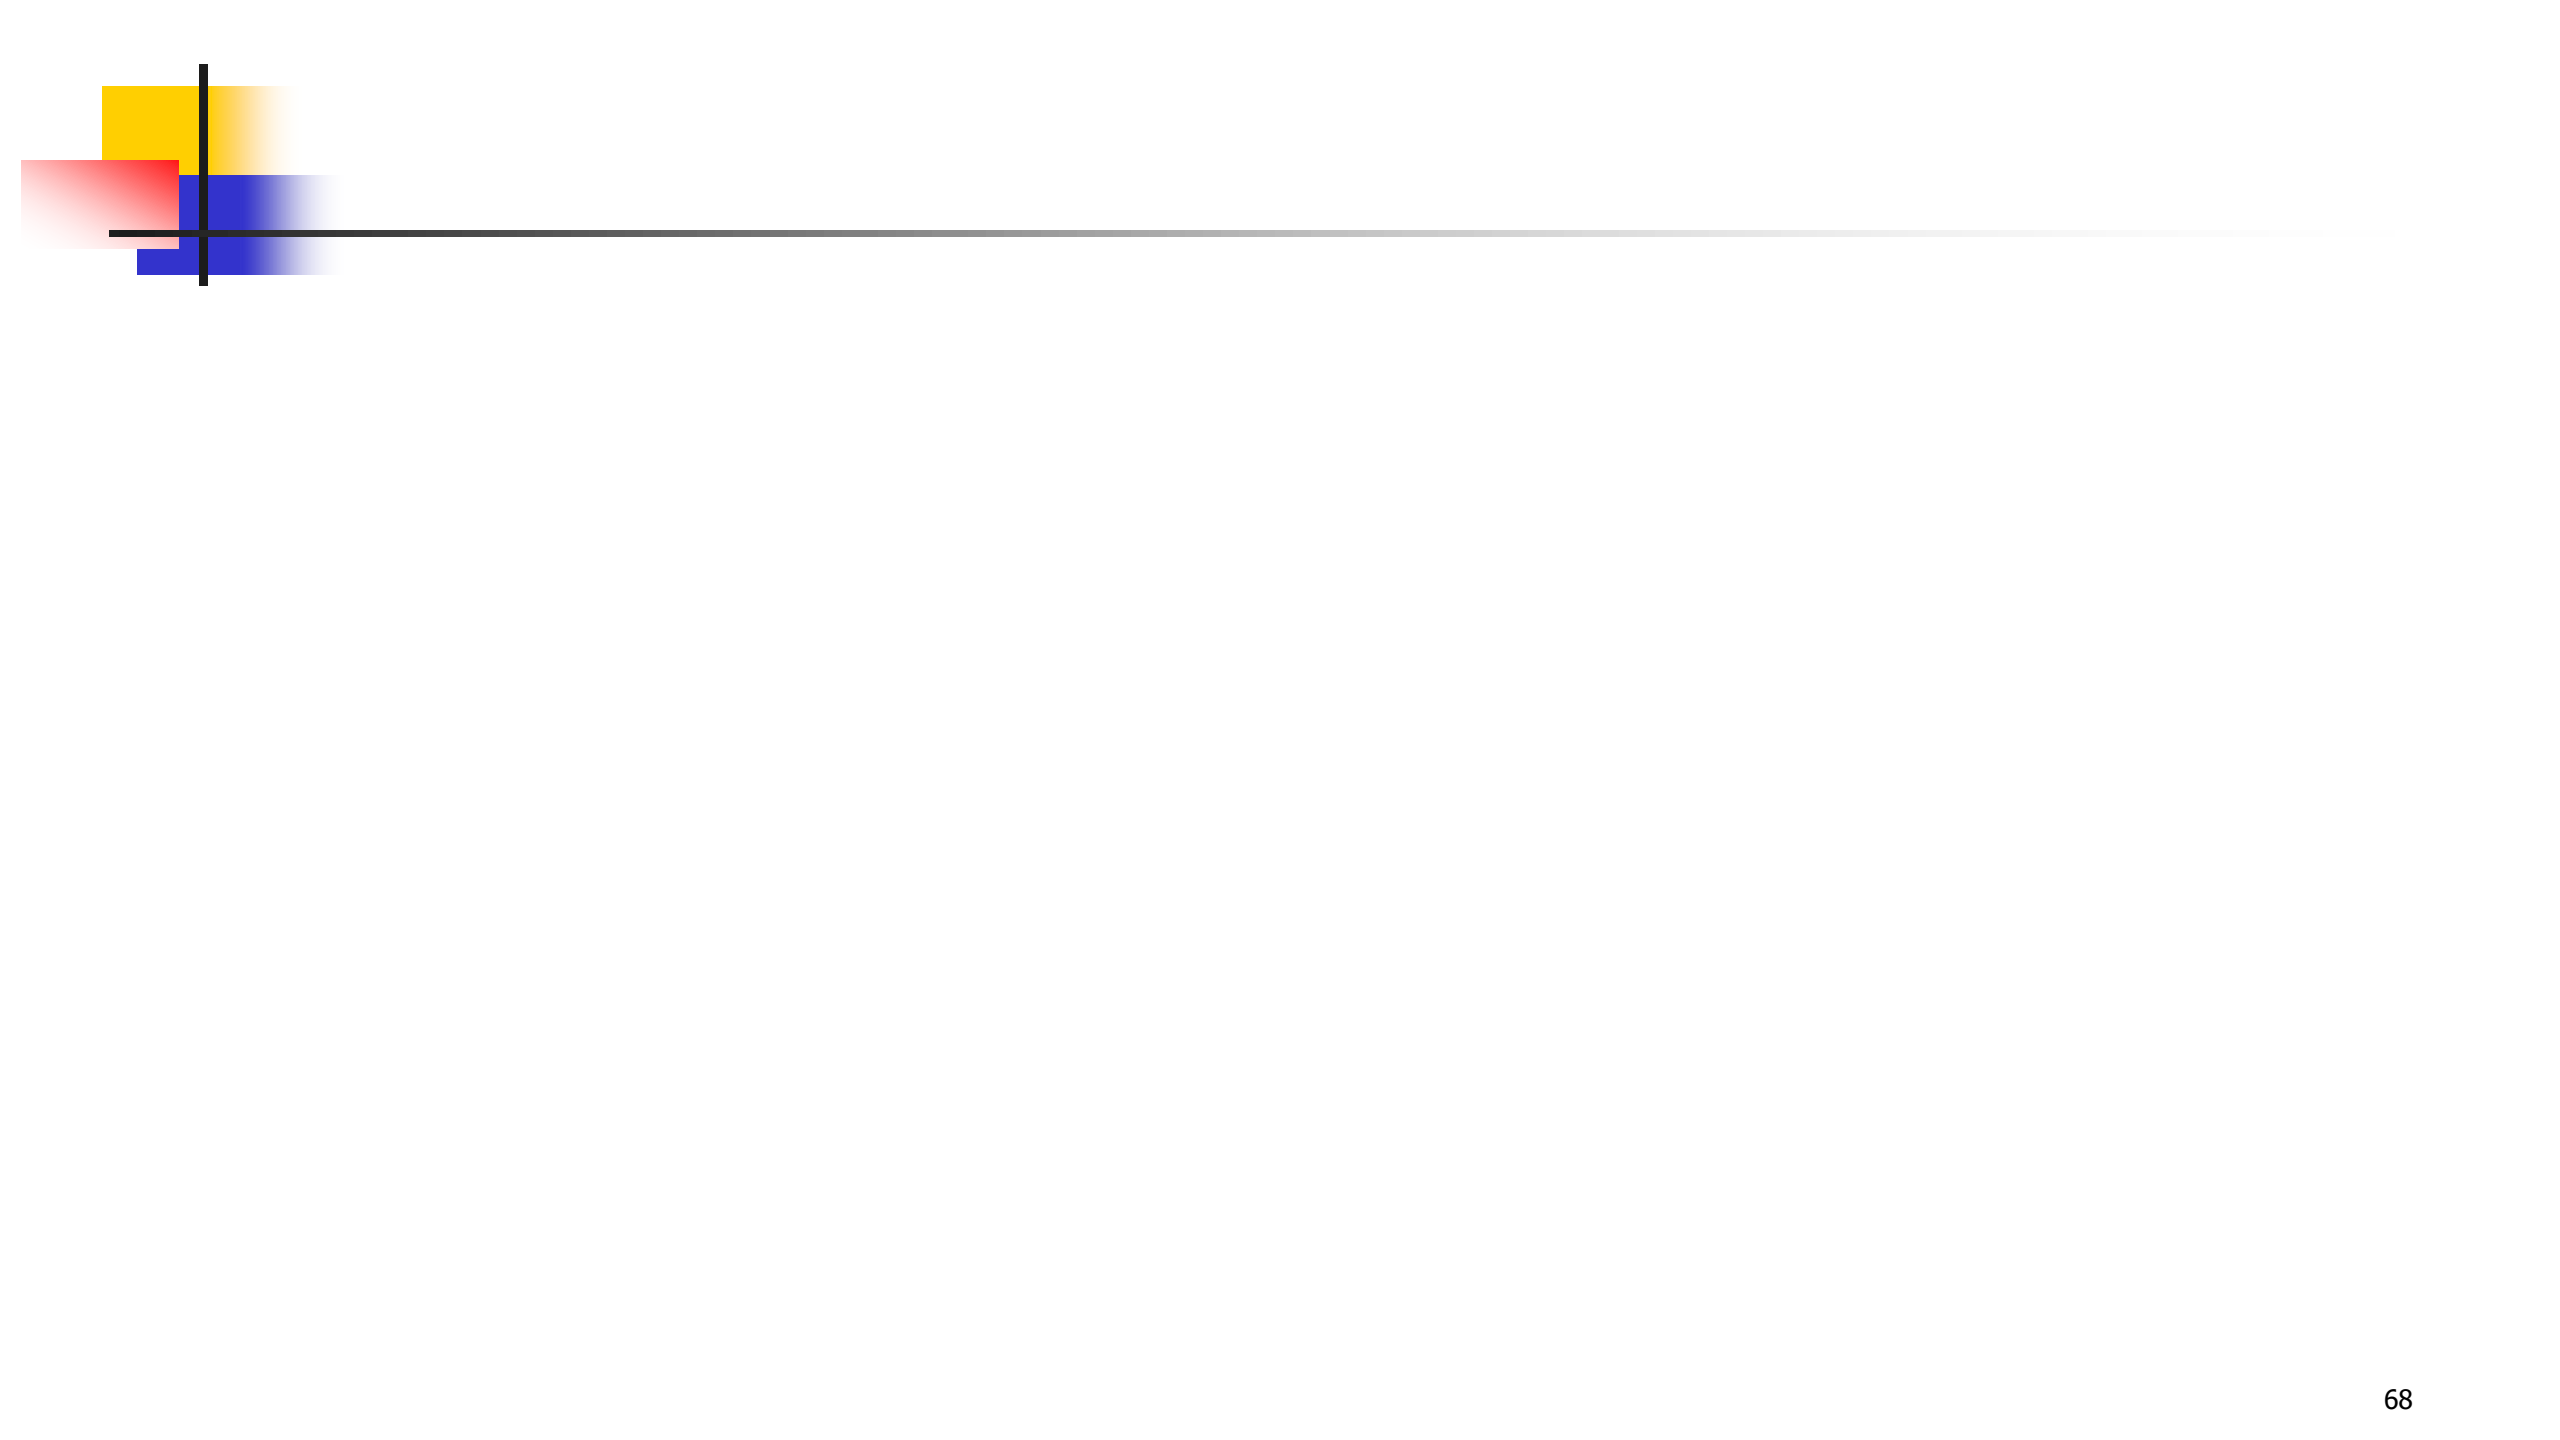
Student

| <u>id</u> | <u>parentId</u> | major | ... |
|-----------|-----------------|-------|-----|
| 1 | 3 | Math | ... |
| ... | ... | ... | ... |
| | | | |



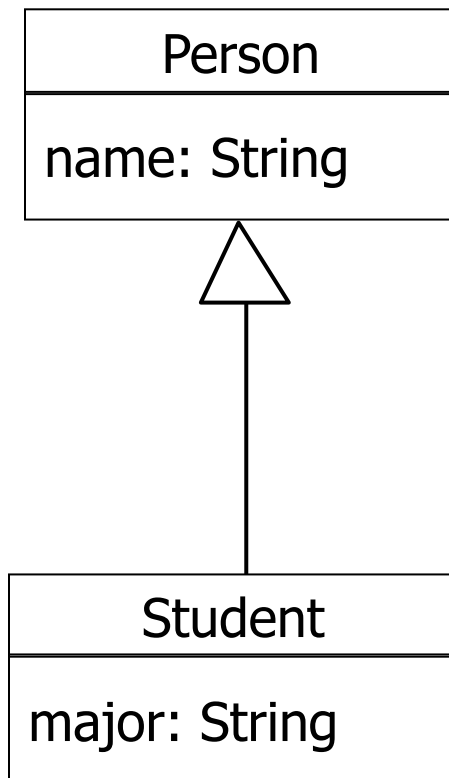
Object persistence

- Method 2:
 - Parent and Child classes are mapped onto a single table, which includes all attributes (from Parent and Child classes)
 - A Parent class object uses only the Parent's columns; the Child columns are wasted
 - A Child class object uses all columns (its columns and the Parent's columns)
 - Mapping a class hierarchy requires creating a table with the union of all attributes and some wasted space is unavoidable



Object persistence

- Generalization mapping, method 2



| PersonHierarchy | | | | Student attrs | |
|-----------------|---------|---------|-----|---------------|------|
| Person attrs | | | | | |
| id | type | name | ... | major | ... |
| 1 | student | William | ... | Math | ... |
| 2 | person | Bill | ... | null | null |
| ... | ... | ... | ... | ... | ... |

`select * from PersonHierarchy where type = 'student'`
to retrieve just Student objects

OR

`select * from PersonHierarchy`
to retrieve both Person and Student objects

- 
-
- See also: <http://agiledata.org/essays/mappingObjects.html>



Object persistence

- A persistence **middleware** system may be used to store and retrieve objects from an RDBMS
- Example: Hibernate (from RedHat)

<http://www.hibernate.org/>

It is a framework for mapping an object-oriented domain model to a relational database

- Mapping is placed in an XML file
- Classes and relationships (1-m, m-n) are mapped onto relational tables
- Objects are stored and retrieved “seamlessly”



Summary: Why are RDBMS useful?

- Data independence – provides abstract view of the data, without details of storage
- Efficient data access – uses techniques to store and retrieve data efficiently
- Reduced application development time – many important functions already supported
- Centralized data administration
- Data Integrity and Security
- Concurrency control and recovery